



Create S3 Buckets with Terraform

K

Kehinde Abiuwa

The screenshot shows the HashiCorp Terraform installation page for macOS. The URL in the browser is <https://developer.hashicorp.com/terraform/install>. The page has a navigation bar with links for Home, Terraform, Install, Tutorials, Documentation, Registry, and Try Cloud. On the left, there's a sidebar with sections for Operating Systems (macOS selected), Windows, Linux, FreeBSD, OpenBSD, Solaris, Release information, and Next steps. The main content area is titled "Install Terraform" and shows options for macOS and Windows. For macOS, it provides package manager instructions ("brew tap hashicorp/tap" and "brew install hashicorp/tap/terraform") and binary download links for AMD64 and ARM64 versions (1.13.0). For Windows, it provides a binary download link. To the right, there's a box titled "About Terraform" with a description and links to "Featured docs" (Introduction to Terraform, Configuration Language, Terraform CLI, HCP Terraform, Provider Use) and a section for "HCP Terraform" with a "Try HCP Terraform for free" button.

Introducing Today's Project!

In this project, I will demonstrate how to install and configure Terraform, set up AWS credentials in the terminal, and use Terraform to create and manage S3 buckets. I will also show how to upload files to S3 with Terraform. The goal is to gain hands-on experience with Terraform for AWS resource management and automation.

Tools and concepts

Services I used were: AWS S3 (bucket creation, static website hosting, object storage). AWS IAM (user and access keys for auth). AWS STS (GetCallerIdentity during credential validation). AWS CLI v2 (to configure creds and verify results). Terraform with the AWS provider. Key concepts I learnt include Infrastructure as Code (IaC) and how Terraform represents it with: Provider, resource, and data blocks (e.g., `aws_s3_bucket`, `aws_s3_object`, `website config`). State management (`terraform.tfstate`) and provider locking (`.terraform.lock.hcl`). The safe workflow: `terraform init` → `plan` → `apply` (prepare, preview, perform). Idempotency & change management (edit config, re-apply to update real infrastructure). S3 public access controls (Public Access Block vs. bucket policies) and implications for website hosting. Credential management (profiles, env vars, and fixing `InvalidClientTokenId` errors). Content deployment via code using `aws_s3_object` to upload files as part of the plan.

Project reflection

This project took me approximately 50 minutes. The most challenging part was debugging the Terraform errors—first the “timeout while waiting for plugin to start” (fixed by re-initializing the provider and clearing macOS quarantine/caches) and then the `InvalidClientTokenId` AWS auth issue (solved by installing AWS CLI, configuring a profile with valid access keys, and verifying with `aws sts get-caller-identity`).



Kehinde Abiuwa
NextWork Student

nextwork.org

It was most rewarding to run terraform apply and see the S3 bucket and website config created, then upload image.png via aws_s3_object and download it from the S3 console to confirm the deployment worked end-to-end.

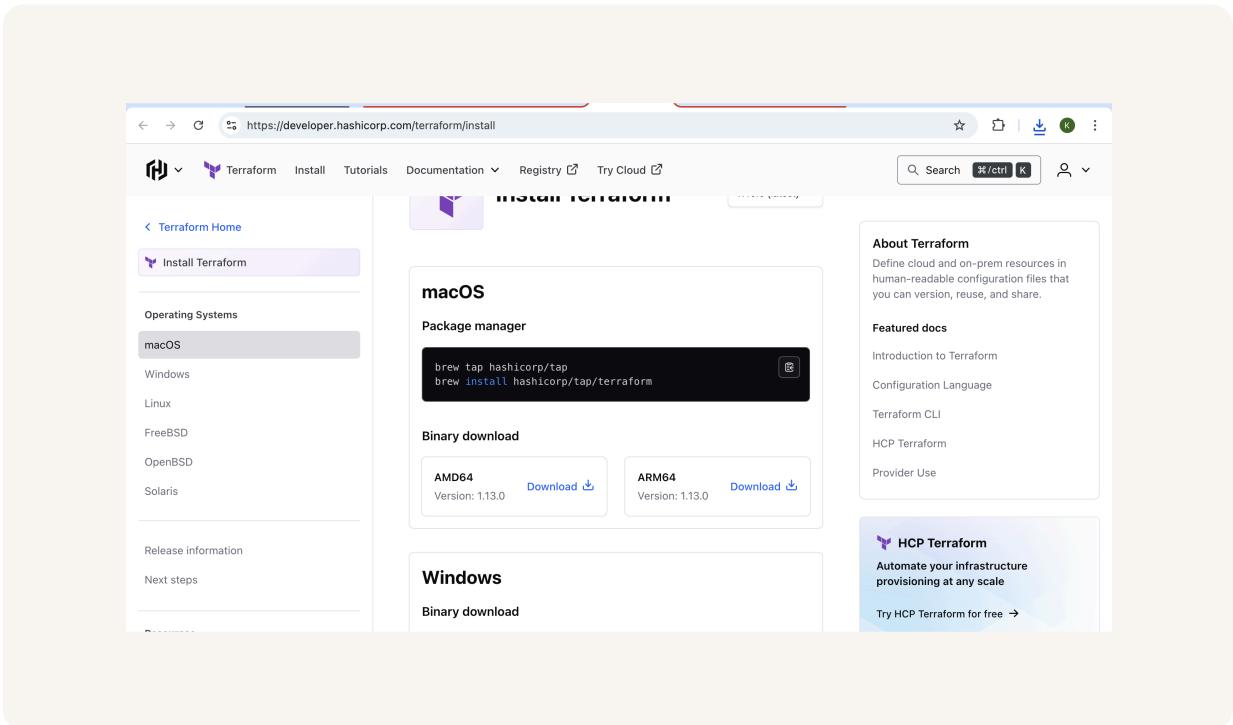
I chose to do this project today because... Something that would make learning with NextWork even better is...

Introducing Terraform

Terraform is an open-source Infrastructure as Code (IaC) tool that lets you define and manage cloud resources (like servers, databases, and storage) using simple configuration files. Instead of manually setting things up in the cloud, you write instructions in Terraform, and it automatically creates, updates, or deletes the resources for you. It works with many providers (like AWS, Azure, Google Cloud), making it easier to automate and manage infrastructure consistently and reliably.

Terraform is one of the most popular tools used for Infrastructure as Code (IaC), which is the practice of managing and provisioning infrastructure through code instead of manual processes. With IaC, you write configuration files that describe the resources you need (like servers, networks, or storage), and tools like Terraform automatically create and manage those resources for you.

Terraform uses configuration files to describe the infrastructure you want to create, change, or manage. These files are written in HashiCorp Configuration Language (HCL) and act as instructions for Terraform. `main.tf` is the primary configuration file where you define the resources (like AWS S3 buckets, EC2 instances, or networking components) that Terraform will build. It usually serves as the starting point of your Terraform project.

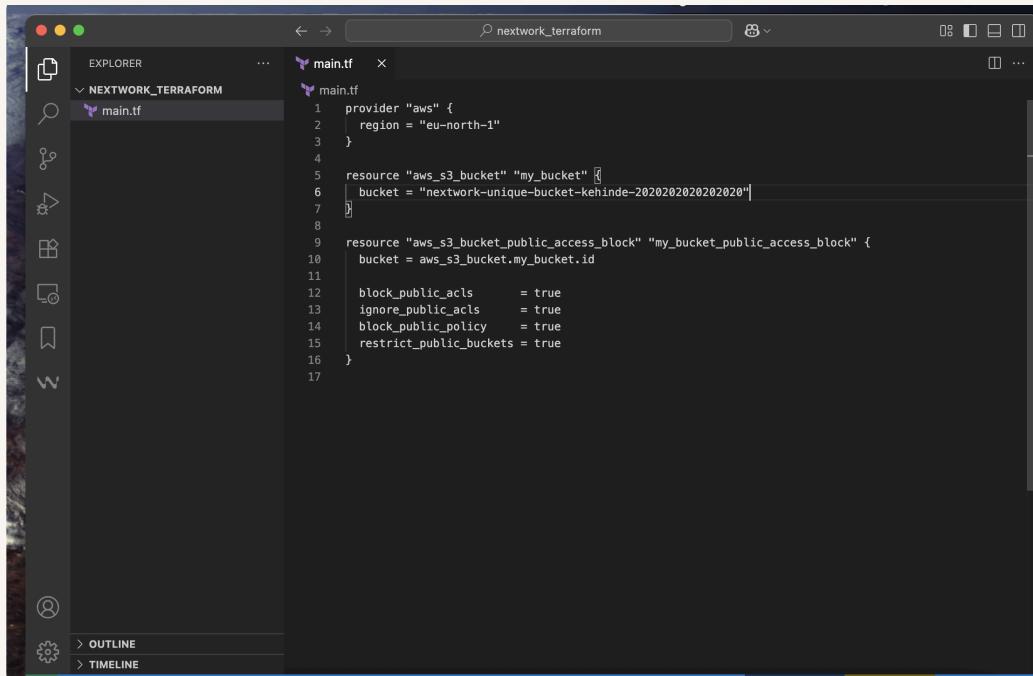


Configuration files

The configuration is structured in blocks. Each block tells Terraform what to do: A provider block defines which cloud provider (AWS) and region to use. A resource block defines what you want to create (like an S3 bucket). Additional resource blocks can manage settings for those resources (like blocking public access to the bucket). The advantage of doing this is that the setup is clear, reusable, and easy to change. Instead of manually creating resources in AWS, you can just edit the configuration file, and Terraform will automatically update your infrastructure to match.

My main.tf configuration has three blocks

The first block indicates the AWS provider and sets the region to eu-north-1, telling Terraform where to create the resources. The second block provisions an Amazon S3 bucket with a unique bucket name. The third block manages the S3 bucket's access settings by blocking all forms of public access, ensuring the bucket stays private and secure.



```
provider "aws" {
  region = "eu-north-1"
}

resource "aws_s3_bucket" "my_bucket" {
  bucket = "nextwork-unique-bucket-kehinde-2020202020202020"
}

resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {
  bucket = aws_s3_bucket.my_bucket.id

  block_public_acls      = true
  ignore_public_acls     = true
  block_public_policy     = true
  restrict_public_buckets = true
}
```

Customizing my S3 Bucket

For my project extension, I visited the official Terraform documentation to explore the `aws_s3_bucket` resource. The documentation shows that Terraform provides a dedicated `aws_s3_bucket` resource to define and manage Amazon S3 buckets declaratively. It includes details on available arguments (such as bucket, acl, versioning, etc.), supported attributes, how to configure additional features like lifecycle rules, server-side encryption, and notes regarding any deprecated options or refactoring updates in newer provider versions

I chose to customise my bucket by enabling S3 static website hosting with an index document (`index.html`), an error document (`error.html`), and a routing rule that redirects any path starting with `docs/` to `documents/`, because I want to serve a simple static site, show a friendly error page, and keep old links working after a folder rename. When I launch my bucket, I can verify my customization by uploading `index.html` and `error.html` and then visiting the bucket's website endpoint (S3 console → Properties → Static website hosting). I should see: the home page loads at the root, a missing page shows my custom error page, a URL like `/docs/example.html` redirects to `/documents/example.html`.



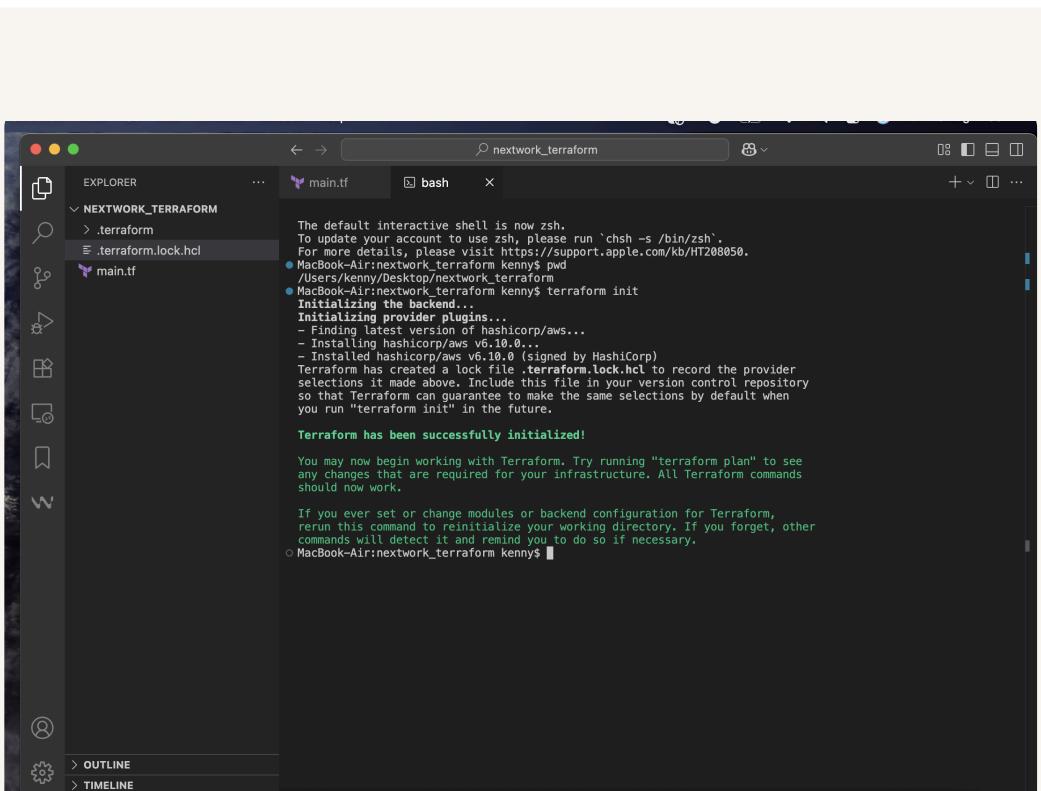
```
EXPLORER      main.tf ×  
NEXTWORK_TERRAFORM  
main.tf  
main.tf  
1 provider "aws" {  
2   region = "eu-north-1"  
3 }  
4  
5 resource "aws_s3_bucket" "my_bucket" {  
6   bucket = "nextwork-unique-bucket-kehinde-20202020202020"  
7 }  
8  
9 resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {  
10   bucket = aws_s3_bucket.my_bucket.id  
11  
12   block_public_acls = true  
13   ignore_public_acls = true  
14   block_public_policy = true  
15   restrict_public_buckets = true  
16 }  
17  
18 resource "aws_s3_bucket_website_configuration" "my_bucket_website_configuration" {  
19   bucket = aws_s3_bucket.my_bucket.id  
20  
21   index_document {  
22     suffix = "index.html"  
23   }  
24  
25   error_document {  
26     key = "error.html"  
27   }  
28  
29   routing_rule {  
30     condition {  
31       key_prefix_equals = "docs/"  
32     }  
33     redirect {  
34       replace_key_prefix_with = "documents/"  
35     }  
36   }  
37 }
```

The screenshot shows a code editor window with a dark theme. The title bar says "nextwork_terraform". The left sidebar has icons for file operations like Open, Save, and Find. The main area shows a Terraform configuration file named "main.tf". The code defines an AWS provider and creates an S3 bucket named "nextwork-unique-bucket-kehinde-20202020202020". It then creates a public access block for the bucket, setting all public access controls to true. Finally, it sets up a website configuration for the bucket,指定 index document为 "index.html"，error document为 "error.html"，并配置了一个路由规则将所有以 "docs/" 开头的请求重定向到 "documents/" 目录。

Terraform commands

I ran `terraform init` to initialize my Terraform project—it sets up the working directory, downloads the required providers (AWS), initializes any modules and the backend (or local state if none is set), and creates the `.terraform/` folder and `.terraform.lock.hcl` for consistent provider versions—so the project is ready to run `terraform plan` and `terraform apply`.

Next, I ran `terraform plan` to safely preview the changes Terraform would make—creating the S3 bucket, applying the website configuration, and setting the Public Access Block



The screenshot shows a Mac OS X terminal window titled "nextwork_terraform". The terminal is displaying the output of the `terraform init` command. The output includes:

```
The default interactive shell is now zsh.  
To update your account to use zsh, please run `chsh -s /bin/zsh`.  
For more details, please visit https://support.apple.com/kb/HT208850.  
● MacBook-Air:nextwork_terraform kenny$ pwd  
/Users/kenny/Desktop/nextwork_terraform  
● MacBook-Air:nextwork_terraform kenny$ terraform init  
Initializing the backend...  
Initializing provider plugins...  
- Found provider/hashicorp/aws v6.10.0  
- Installing hashicorp/aws v6.10.0  
- Installed hashicorp/aws v6.10.0 (signed by HashiCorp)  
Terraform has created a lock file .terraform.lock.hcl to record the provider  
selections it made above. Include this file in your version control repository  
so that Terraform can guarantee to make the same selections by default when  
you run "terraform init" in the future.  
Terraform has been successfully initialized!  
You may now begin working with Terraform. Try running "terraform plan" to see  
any changes that are required for your infrastructure. All Terraform commands  
should now work.  
If you ever set or change modules or backend configuration for Terraform,  
rerun this command to reinitialize your working directory. If you forget, other  
commands will detect it and remind you to do so if necessary.  
○ MacBook-Air:nextwork_terraform kenny$
```

AWS CLI and Access Keys

When I tried to plan my Terraform configuration, I received an error message that says "InvalidClientTokenId: The security token included in the request is invalid" because Terraform called AWS STS to identify my account, but my AWS credentials/profile were incorrect or expired (e.g., wrong access keys or a stale AWS_SESSION_TOKEN), so authentication failed.

To resolve my error, first I installed AWS CLI, which is Amazon's command-line tool for authenticating to AWS and managing services. It provides the aws command and stores credentials in `~/.aws/credentials` and `~/.aws/config`—the same locations Terraform reads—so it's the easiest way to set up valid creds.

I set up AWS access keys to securely authenticate the AWS CLI and Terraform to my AWS account, so they can sign API requests and create/manage resources during plan and apply without using the web console.



```
Planning failed. Terraform encountered an error while generating this plan.

| Error: Retrieving AWS account details: validating provider credentials: retrieving caller identity from STS: operation error STS: GetCallerIdentity, https response error StatusCode: 403, RequestID: ab0d4481-f21d-4052-b8e5-4795ebcb5e2f, api error InvalidClientTokenId: The security token included in the request is invalid.

with provider["registry.terraform.io/hashicorp/aws"],
on main.tf line 1, in provider "aws":
  1: provider "aws" {

2025-08-26T19:22:48.615+0100 [DEBUG] provider.stdio: received EOF, stopping recv loop: err="rpc error: code = Unavailable desc = error reading from server: EOF"
2025-08-26T19:22:48.627+0100 [INFO] provider: plugin process exited: plugin=terraform/providers/registry.terraform.io/hashicorp/aws/6.10.0/darwin_amd64/terraform-provider-aws_v6.10.0_x5_id=91390
2025-08-26T19:22:48.627+0100 [DEBUG] provider: plugin exited
MacBook-Air:nextwork_terraform kenny$ terraform plan
| Error: timeout while waiting for plugin to start

● MacBook-Air:nextwork_terraform kenny$ clear
● MacBook-Air:nextwork_terraform kenny$ aws --version
aws-cli/2.28.15 Python/3.13.4 Darwin/24.6.0 exe/x86_64
○ MacBook-Air:nextwork_terraform kenny$ █
```

> OUTLINE
> TIMELINE

Finish Setup Windsurf: (...) □



Lanching the S3 Bucket

I ran `terraform apply` to execute the plan and create the real AWS resources defined in `main.tf`—the S3 bucket, its Public Access Block, and the website configuration—after previewing the changes with `terraform plan`. Running `terraform apply` will affect my AWS account by: Creating/modifying/deleting resources to match my config (in this case: creating the S3 bucket and its settings). Updating the Terraform state file (`terraform.tfstate`) with the real IDs/ARNs so future runs know what exists. Incurring charges for any billable resources and data usage (S3 storage, requests, and data transfer). Applying access controls I defined (e.g., blocking public access or enabling website hosting), which changes who can access the bucket. Recording a change history in AWS (CloudTrail shows API calls made by Terraform's AWS provider).

The sequence is crucial because each command prepares the next one and protects you from accidental changes. `terraform init` → `setup`. Downloads/locks the provider(s), configures the backend/state, and makes the working directory usable. Run it first, and again whenever you add/upgrade providers or change the backend. `terraform plan` → `preview`. Reads your `.tf` files and current state, compares to real AWS, and shows the exact actions (create/change/destroy) without changing anything. Use it to verify intent. Optionally save it: `terraform plan -out=tfplan`. `terraform apply` → `execute`. Performs the actions to match your config and updates `terraform.tfstate`. Best practice: apply the saved plan (`terraform apply tfplan`) so you execute exactly what you reviewed.



The screenshot shows the AWS S3 Buckets page. The top navigation bar includes the AWS logo, search bar, and account information (Account ID: 5112-3943-1868, Europe (Stockholm)). The left sidebar lists various S3 features: General purpose buckets, Directory buckets, Table buckets, Vector buckets, Access Grants, Access Points (General Purpose Buckets, FSx file systems), Access Points (Directory Buckets), Object Lambda Access Points, Multi-Region Access Points, Batch Operations, and IAM Access Analyzer for S3. A note at the bottom of the sidebar says "Block Public Access settings for this account". The main content area is titled "General purpose buckets (1)" and shows a table with one row. The table columns are Name, AWS Region, and Creation date. The single entry is "nextwork-unique-bucket-kehinde-20202020202020", located in "Europe (Stockholm) eu-north-1" and created on "August 26, 2025, 20:00:50 (UTC+01:00)". There are buttons for Copy ARN, Empty, Delete, and Create bucket. Below the table are two cards: "Account snapshot" (Updated daily) which provides visibility into storage usage and activity trends, and "External access summary - new" (Updated daily) which helps identify bucket permissions allowing public access or access from other AWS accounts.

Name	AWS Region	Creation date
nextwork-unique-bucket-kehinde-20202020202020	Europe (Stockholm) eu-north-1	August 26, 2025, 20:00:50 (UTC+01:00)



Uploading an S3 Object

I created a new resource block to upload a local file to my S3 bucket using:

```
resource "aws_s3_object" "image" { bucket = aws_s3_bucket.my_bucket.id key = "image.png" # path/name in the bucket source = "image.png" # local file to upload # optional: # content_type = "image/png" }
```

because I want Terraform to manage my content as code—each apply ensures the file exists in the bucket with the expected name and content, making deployments repeatable.

We need to run `terraform apply` again because we changed the configuration (added a new `aws_s3_object` to upload `image.png`). Terraform only makes real changes when you apply, so running it again will:

- Create the new object in the S3 bucket (and any other updates you added).
- Sync state so `terraform.tfstate` reflects the new resources.

Best practice: run `terraform plan` first to preview the delta, then `terraform apply` to push it.

To validate that I've updated my configuration successfully, I opened the S3 console in my browser, navigated to my bucket, downloaded `image.png`, and confirmed it was the same image I had just uploaded—verifying the object upload worked as intended.



The screenshot shows a code editor window with a dark theme. The left sidebar is labeled 'EXPLORER' and shows a folder structure for 'NEXTWORK_TERRAFORM' containing '.terraform', '.terraform.lock.hcl', 'image.png', 'main.tf', and 'terraform.tfstate'. The main pane displays the contents of 'main.tf'.

```
provider "aws" {
  region = "eu-north-1"
}

resource "aws_s3_bucket" "my_bucket" {
  bucket = "nextwork-unique-bucket-kehinde-2020202020202020"
}

resource "aws_s3_bucket_public_access_block" "my_bucket_public_access_block" {
  bucket = aws_s3_bucket.my_bucket.id

  block_public_acls      = true
  ignore_public_acls     = true
  block_public_policy    = true
  restrict_public_buckets = true
}

resource "aws_s3_object" "image" {
  bucket = aws_s3_bucket.my_bucket.id # Reference the bucket ID
  key    = "image.png" # Path in the bucket
  source = "image.png" # Local file path
}
```



nextwork.org

The place to learn & showcase your skills

Check out nextwork.org for more projects

