



Launch a Kubernetes Cluster

K

Kehinde Abiuwa

The screenshot shows the AWS EKS console interface. The top navigation bar includes the AWS logo, search bar, account ID (5112-3943-1868), region (Europe (Stockholm)), and IAM role (kehindeabiuwa-IAM-Admin). The main navigation menu on the left lists 'Amazon Elastic Kubernetes Service' (selected), 'Clusters' (highlighted in blue), 'Settings' (with 'Dashboard settings' and 'Console settings' options), 'Amazon EKS Anywhere' (with 'Enterprise Subscriptions' option), 'Related services' (with 'Amazon ECR' and 'AWS Batch' options), and 'Documentation'. The central content area displays a success message: 'Access entry was successfully created.' Below this is a table titled 'Nodes (3) info' showing three nodes:

Node name	Instance type	Compute	Managed by	Created	Status
ip-192-168-28-6.eu-north-1.compute.internal	t3.micro	Node group	nextwork-nodegroup	27 minutes ago	Read
ip-192-168-63-183.eu-north-1.compute.internal	t3.micro	Node group	nextwork-nodegroup	27 minutes ago	Read
ip-192-168-83-47.eu-north-1.compute.internal	t3.micro	Node group	nextwork-nodegroup	27 minutes ago	Read

Below the nodes is a section titled 'Node groups (1) info' with one item listed:

Group name	Desired size	AMI release version	Launch template
nextwork-nodegroup	3	v3.1.1	AmazonLinux2



Introducing Today's Project!

In this project, I will launch and connect to an EC2 instance, create a Kubernetes cluster, monitor its creation with CloudFormation, access it via an IAM access entry, and run a resilience test because I want hands-on mastery of building, securing, observing, and hardening cloud-native infrastructure on AWS.



What is Kubernetes?

Kubernetes is an open-source system that runs lots of containerized apps across many machines and automates the boring-but-critical stuff: where apps run, how they talk, how they scale, how they heal, and how they're updated—without you babysitting each server. Companies and developers use Kubernetes to:

- Keep apps reliable
- Scale automatically
- Ship faster, safer
- Run anywhere
- Use hardware efficiently
- Organize microservices
- Manage config & secrets
- Standardize platforms
- Secure multi-tenant envs
- Extend operations

I used eksctl to spin up an Amazon EKS cluster from my EC2 box. eksctl took my flags, built CloudFormation stacks, and created the EKS control plane, a managed node group, default VPC networking (since I didn't pass a custom VPC), IAM roles, and a kubeconfig so kubectl can talk to the cluster. The create cluster command I ran defined:

```
Cluster name: nextwork-eks-cluster
Region: eu-north-1
Kubernetes version: 1.33
Node group name: nextwork-nodegroup
Instance type: t3.micro
Capacity: desired 3 nodes, with autoscaling min 1 and max 3
```

I initially ran into two errors while using eksctl. The first one was because eksctl wasn't installed/on my PATH, so the shell returned eksctl: command not found. I fixed it by downloading the binary, moving it to /usr/local/bin, and verifying with eksctl version. The second one was because my EC2 instance didn't have AWS credentials/permissions (no IAM instance role and/or region set). eksctl couldn't call EKS/CloudFormation/EC2 and would show "Unable to locate credentials" or "AccessDenied." I fixed it by attaching an appropriate IAM role (with EKS/CFN/EC2 + IAM PassRole rights) to the instance, setting the region, confirming with aws sts get-caller-identity, and re-running the command.



```
# Amazon Linux 2023
--> https://aws.amazon.com/linux/amazon-linux-2023

[ec2-user@ip-172-31-35-98 ~]$ eksctl create cluster \
--name nextwork-eks-cluster \
--nodegroup-name nextwork-nodegroup \
--node-type t3.micro \
--nodes-min 1 \
--nodes-max 3 \
--version 1.33
sh: eksctl: Command not found
[ec2-user@ip-172-31-35-98 ~]$
```



eksctl and CloudFormation

CloudFormation helped create my EKS cluster because eksctl builds everything through CloudFormation stacks. Those stacks declare all the AWS pieces the cluster needs—then CloudFormation creates them in order, tracks progress in the Events tab, and rolls back safely if something fails. Concretely, it stood up the EKS control plane, security groups, IAM roles/policies (cluster + node role), and the managed node group (with its launch template and autoscaling). It created VPC resources because I didn't provide an existing VPC/subnets. By default, eksctl asks CloudFormation to make a fresh VPC with public and private subnets across multiple AZs, route tables, an Internet Gateway, NAT Gateway + EIP (so private nodes can reach the internet for images/updates), and the necessary security group rules—giving the cluster sane, production-style networking out of the box.

There was also a second CloudFormation stack for the managed node group (nextwork-nodegroup). eksctl always makes two stacks: one for the cluster/control plane and another for each node group. The node-group stack created the EC2 capacity (Auto Scaling group + Launch Template), the NodeInstanceRole (IAM), and the node security group/bootstrapping so the nodes can join the cluster. The difference between a cluster and node group is: Cluster → the EKS control plane (API server, etcd, cluster security groups/roles, endpoint). It can exist by itself but can't run Pods without workers. Node group → a set of worker nodes (EC2 instances) that register to the cluster and run your Pods; you can have many, each with its own instance types/size/spot settings, and you can scale/upgrade/delete them independently of the cluster.

The screenshot shows the AWS CloudFormation console interface. On the left, there's a navigation sidebar with sections like CloudFormation, Infrastructure Composer, Hooks, and Registry. The main area displays two active stacks:

- eksctl-nextwork-eks-cluster-nodegroup-nextwork-nodegroup**: Status: CREATE_COMPLETE. Last updated: 2025-08-27 01:07:30 UTC+0100.
- eksctl-nextwork-eks-cluster-cluster**: Status: CREATE_COMPLETE. Last updated: 2025-08-27 00:57:26 UTC+0100.

Below the stacks, there's a table of events:

Timestamp	Logical ID	Status
2025-08-27 01:09:59 UTC+0100	eksctl-nextwork-eks-cluster-nodegroup-nextwork-nodegroup	CREATE_COMPLETE
2025-08-27 01:09:58 UTC+0100	ManagedNodeGroup	CREATE_COMPLETE
2025-08-27 01:07:55 UTC+0100	ManagedNodeGroup	CREATE_IN_PROGRESS
2025-08-27 01:07:52 UTC+0100	ManagedNodeGroup	CREATE_IN_PROGRESS
2025-08-27 01:07:51 UTC+0100	NodeInstanceRole	CREATE_COMPLETE
2025-08-27 01:07:34	LaunchTemplate	CREATE_COMPLETE

The EKS console

I had to create an IAM access entry in order to map my AWS identity (the EC2 instance role) to Kubernetes RBAC, so the cluster would actually authorize my kubectl calls (see nodes, deploy apps, etc.). AWS admin rights don't automatically carry into Kubernetes—that's why the console warned me. An access entry is an EKS-managed link between an IAM user/role and a set of EKS access policies (e.g., cluster-admin, read-only, namespace-scoped). EKS then creates/updates the underlying Kubernetes RBAC and aws-auth mapping for you. I set it up by opening EKS Console → Cluster → Access → Grant access, selecting the IAM role my EC2 uses, choosing a policy (lab: AmazonEKSClusterAdminPolicy; prod: least-privilege and/or namespace scope), and saving. Then I refreshed my EKS cluster console page and I could then see the nodes

It took about 20 minutes to create my cluster. Since I'll create this cluster again in the next project of this series, maybe this process could be sped up if I reuse the same cluster (or an existing VPC), start with a single node or Fargate, and use a supported version—so CloudFormation has less to build and the nodes join faster



The screenshot shows the AWS EKS console interface. The top navigation bar includes the AWS logo, search bar, and account information (Account ID: 5112-3943-1868, Europe (Stockholm), kehindeabiuwa-IAM-Admin). The left sidebar has sections for Dashboard (New), Clusters, Settings (Dashboard settings, Console settings), Amazon EKS Anywhere (Enterprise Subscriptions), Related services (Amazon ECR, AWS Batch), and Documentation. The main content area shows a green notification bar stating "Access entry was successfully created." Below it is a "Nodes (3) Info" section with a table:

Node name	Instance type	Compute	Managed by	Created	Status
ip-192-168-28-6.eu-north-1.compute.internal	t3.micro	Node group	nextwork-nodegroup	27 minutes ago	Read
ip-192-168-63-183.eu-north-1.compute.internal	t3.micro	Node group	nextwork-nodegroup	27 minutes ago	Read
ip-192-168-83-47.eu-north-1.compute.internal	t3.micro	Node group	nextwork-nodegroup	27 minutes ago	Read

Below the nodes table is a "Node groups (1) Info" section with a table:

Group name	Desired size	AMI release version	Launch template
nextwork-nodegroup	1	2023.09.14.0.0.0	AmazonLinux2



EXTRA: Deleting nodes

Did you know you can find your EKS cluster's nodes in Amazon EC2? This is because when you use managed (or self-managed) node groups, the "nodes" are just regular EC2 instances running the EKS-optimized AMI (or Bottlerocket) in your VPC. EKS manages the control plane for you, but it tells Auto Scaling to launch worker EC2s in your account—so they appear in the EC2 console like any other instance.

Desired size means the current target number of nodes in your EKS managed node group/Auto Scaling group. The ASG keeps launching or terminating EC2s to stay at this number (so if you kill a node, it replaces it to get back to the desired count). Minimum and maximum sizes are helpful for setting the floor and ceiling the group can scale between: Minimum = the baseline you always want running (e.g., enough nodes for HA across AZs and to fit DaemonSets/system pods). Maximum = the cap to prevent runaway scale and control costs/quotas.

When I deleted my EC2 instances, Kubernetes marked those nodes NotReady and the Auto Scaling Group for the managed node group launched replacement EC2 instances to bring the node count back to the desired size. Once the new instances booted and joined the cluster, everything returned to steady state. This is because EKS worker nodes are just EC2 instances managed by an Auto Scaling Group, and Kubernetes controllers keep the pod count at the desired state. If there isn't enough capacity while replacements are starting, some pods will sit Pending briefly.



The screenshot shows the AWS EC2 Instances page with the following details:

Instances (7) Info

Last updated less than a minute ago

Actions ▾ Launch instances ▾

Find Instance by attribute or tag (case-sensitive)

All states ▾

Name	Instance ID	Instance state	Instance type	Status check	Alarm status
nextwork-eks-...	i-0021ac85b1faf5ef3	Terminated	t3.micro	-	View alarms +
nextwork-eks-...	i-0c4b1c4746db14aa3	Running	t3.micro	3/3 checks passed	View alarms +
nextwork-eks-...	i-0aa285a5fc9d69533	Running	t3.micro	3/3 checks passed	View alarms +
nextwork-eks-...	i-0555401997d070a6c	Terminated	t3.micro	-	View alarms +
nextwork-eks-...	i-0ce2b60f2d9355d67	Running	t3.micro	3/3 checks passed	View alarms +
nextwork-eks-...	i-02629283e6e83c2fa	Terminated	t3.micro	-	View alarms +
nextwork-eks-...	i-0286b9ac47335725e	Running	t3.micro	Initializing	View alarms +

Select an instance



nextwork.org

The place to learn & showcase your skills

Check out nextwork.org for more projects

